



ROCm

Release Notes v2.10

Publication #	1.0	Revision:	2.1
Issue Date:	November 2019		

Specification Agreement

This Specification Agreement (this “Agreement”) is a legal agreement between Advanced Micro Devices, Inc. (“AMD”) and “You” as the recipient of the attached AMD Specification (the “Specification”). If you are accessing the Specification as part of your performance of work for another party, you acknowledge that you have authority to bind such party to the terms and conditions of this Agreement. If you accessed the Specification by any means or otherwise use or provide Feedback (defined below) on the Specification, You agree to the terms and conditions set forth in this Agreement. If You do not agree to the terms and conditions set forth in this Agreement, you are not licensed to use the Specification; do not use, access or provide Feedback about the Specification. In consideration of Your use or access of the Specification (in whole or in part), the receipt and sufficiency of which are acknowledged, You agree as follows:

1. You may review the Specification only (a) as a reference to assist You in planning and designing Your product, service or technology (“Product”) to interface with an AMD product in compliance with the requirements as set forth in the Specification and (b) to provide Feedback about the information disclosed in the Specification to AMD.
2. Except as expressly set forth in Paragraph 1, all rights in and to the Specification are retained by AMD. This Agreement does not give You any rights under any AMD patents, copyrights, trademarks or other intellectual property rights. You may not (i) duplicate any part of the Specification; (ii) remove this Agreement or any notices from the Specification, or (iii) give any part of the Specification, or assign or otherwise provide Your rights under this Agreement, to anyone else.
3. The Specification may contain preliminary information, errors, or inaccuracies, or may not include certain necessary information. Additionally, AMD reserves the right to discontinue or make changes to the Specification and its products at any time without notice. The Specification is provided entirely “AS IS.” AMD MAKES NO WARRANTY OF ANY KIND AND DISCLAIMS ALL EXPRESS, IMPLIED AND STATUTORY WARRANTIES, INCLUDING BUT NOT LIMITED TO IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, NONINFRINGEMENT, TITLE OR THOSE WARRANTIES ARISING AS A COURSE OF DEALING OR CUSTOM OF TRADE. AMD SHALL NOT BE LIABLE FOR DIRECT, INDIRECT, CONSEQUENTIAL, SPECIAL, INCIDENTAL, PUNITIVE OR EXEMPLARY DAMAGES OF ANY KIND (INCLUDING LOSS OF BUSINESS, LOSS OF INFORMATION OR DATA, LOST PROFITS, LOSS OF CAPITAL, LOSS OF GOODWILL) REGARDLESS OF THE FORM OF ACTION WHETHER IN CONTRACT, TORT (INCLUDING NEGLIGENCE) AND STRICT PRODUCT LIABILITY OR OTHERWISE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGES.
4. Furthermore, AMD’s products are not designed, intended, authorized or warranted for use as components in systems intended for surgical implant into the body, or in other applications intended to support or sustain life, or in any other application in which the failure of AMD’s product could create a situation where personal injury, death, or severe property or environmental damage may occur.
5. You have no obligation to give AMD any suggestions, comments or feedback (“Feedback”) relating to the Specification. However, any Feedback You voluntarily provide may be used by AMD without restriction, fee or obligation of confidentiality. Accordingly, if You do give AMD Feedback on any version of the Specification, You agree AMD may freely use, reproduce, license, distribute, and otherwise commercialize Your Feedback in any product, as well as has the right to sublicense third parties to do the same. Further, You will not give AMD any Feedback that You may have reason to believe is (i) subject to any patent, copyright or other intellectual property claim or right of any third party; or (ii) subject to license terms which seek to require any product or intellectual property incorporating or derived from Feedback or any Product or other AMD intellectual property to be licensed to or otherwise provided to any third party.
6. You shall adhere to all applicable U.S. import/export laws and regulations, as well as the import/export control laws and regulations of other countries as applicable. You further agree to not export, re-export, or transfer, directly or indirectly, any product, technical data, software or source code received from AMD under this license, or the direct product of such technical data or software to any country for which the United States or any other applicable government requires an export license or other governmental approval without first obtaining such licenses or approvals; or in violation of any applicable laws or regulations of the United States or the country where the technical

data or software was obtained. You acknowledge the technical data and software received will not, in the absence of authorization from U.S. or local law and regulations as applicable, be used by or exported, re-exported or transferred to: (i) any sanctioned or embargoed country, or to nationals or residents of such countries; (ii) any restricted end-user as identified on any applicable government end-user list; or (iii) any party where the end-use involves nuclear, chemical/biological weapons, rocket systems, or unmanned air vehicles. For the most current Country Group listings, or for additional information about the EAR or Your obligations under those regulations, please refer to the U.S. Bureau of Industry and Security's website at <http://www.bis.doc.gov/>.

7. The Software and related documentation are "commercial items", as that term is defined at 48 C.F.R. §2.101, consisting of "commercial computer software" and "commercial computer software documentation", as such terms are used in 48 C.F.R. §12.212 and 48 C.F.R. §227.7202, respectively. Consistent with 48 C.F.R. §12.212 or 48 C.F.R. §227.7202-1 through 227.7202-4, as applicable, the commercial computer software and commercial computer software documentation are being licensed to U.S. Government end users (a) only as commercial items and (b) with only those rights as are granted to all other end users pursuant to the terms and conditions set forth in this Agreement. Unpublished rights are reserved under the copyright laws of the United States.

8. This Agreement is governed by the laws of the State of California without regard to its choice of law principles. Any dispute involving it must be brought in a court having jurisdiction of such dispute in Santa Clara County, California, and You waive any defenses and rights allowing the dispute to be litigated elsewhere. If any part of this agreement is unenforceable, it will be considered modified to the extent necessary to make it enforceable, and the remainder shall continue in effect. The failure of AMD to enforce any rights granted hereunder or to take action against You in the event of any breach hereunder shall not be deemed a waiver by AMD as to subsequent enforcement of rights or subsequent actions in the event of future breaches. This Agreement is the entire agreement between You and AMD concerning the Specification; it may be changed only by a written document signed by both You and an authorized representative of AMD.

DISCLAIMER

The information contained herein is for informational purposes only and is subject to change without notice. While every precaution has been taken in the preparation of this document, it may contain technical inaccuracies, omissions and typographical errors, and AMD is under no obligation to update or otherwise correct this information. Advanced Micro Devices, Inc. makes no representations or warranties with respect to the accuracy or completeness of the contents of this document, and assumes no liability of any kind, including the implied warranties of noninfringement, merchantability or fitness for particular purposes, with respect to the operation or use of AMD hardware, software or other products described herein. No license, including implied or arising by estoppel, to any intellectual property rights is granted by this document. Terms and limitations applicable to the purchase or use of AMD's products are as set forth in a signed agreement between the parties or in AMD's Standard Terms and Conditions of Sale. S

AMD, the AMD Arrow logo, Radeon, Ryzen, Epyc, and combinations thereof are trademarks of Advanced Micro Devices, Inc.

Google® is a registered trademark of Google LLC.

PCIe® is a registered trademark of PCI-SIG Corporation.

Linux is the registered trademark of Linus Torvalds in the U.S. and other countries.

Ubuntu and the Ubuntu logo are registered trademarks of Canonical Ltd.

Other product names used in this publication are for identification purposes only and may be trademarks of their respective companies.

As of Jun 19, 2019, Radeon Instinct™ MI50 and MI60 "Vega 7nm" technology-based accelerators support PCIe® Gen 4.0* providing up to 64 GB/s peak theoretical transport data bandwidth from CPL to GPU per card. Previous Gen Radeon Instinct compute GPU cards are based on PCIe Gen 3.0 providing up to 32 GB/s peak theoretical transport rate bandwidth performance. Peak theoretical transport rate performance is calculated by Baud Rate * width in bytes * # directions = GB/s per card. PCIe Gen3: 8 * 2 * 2 = 32 GB/s. PCIe Gen4: 16 * 2 * 2 = 64 GB/s. Radeon Instinct™ MI50 and MI60 "Vega 7nm" technology-based accelerators include dual Infinity Fabric™ Links providing up to 184 GB/s peak theoretical GPU to GPU or Peer-to-Peer (P2P) transport rate bandwidth performance per GPU card. Combined with PCIe Gen 4 compatibility providing an aggregate GPU card I/O peak bandwidth of up to 248 GB/s. Performance guidelines are estimated only and may vary. Previous Gen Radeon Instinct compute GPU cards provide up to 32 GB/s peak PCIe Gen 3.0 bandwidth performance. Infinity Fabric Link technology peak theoretical transport rate performance is calculated by Baud Rate * width in bytes * # directions * # links = GB/s per card. Infinity Fabric Link: 23 * 2 * 2 = 92 GB/s. MI50 | MI60 each have two links: 92 GB/s * 2 links per GPU = 184 GB/s. Refer to server manufacture PCIe Gen 4.0 compatibility and performance guidelines for potential peak performance of the specified server model numbers. Server manufacturers may vary configuration offerings yielding different results. <https://pcisig.com/>, <https://www.chipestimate.com/PCI-Ex-press-Gen-4-a-Big-Pipe-for-Big-Data/Cadence/Technical-Article/2014/04/15>, <https://www.tomshardware.com/news/pci-4.0-power-speed-express,32525.html> AMD has not independently tested or verified external/third party results/data and bears no responsibility for any errors or omissions therein. R19-1A

Table of Contents

Table of Contents	5
Revision History	7
Chapter 1 Introduction.....	8
1.1 What Is ROCm?.....	8
1.1.1 ROCm Components.....	8
1.1.2 Supported Operating Systems.....	9
1.1.3 Important ROCm Links	9
Chapter 2 What's New in Version 2.10.....	10
2.1 New Features	10
2.1.1 rocBLAS - Support for Complex GEMM	10
2.1.2 Support for SLES 15 SP1	10
2.1.3 Code Marker Support for rocProfiler and rocTracer Libraries	12
Chapter 3 Fixed Issues.....	12
3.1 Fixed Issues in the v2.10 Release	12
3.1.1 Running TensorFlow and PyTorch Frameworks Consecutively Results in the Memory Access Fault error	12
3.1.2 Issue with the Docker Container Environment Variable Setting.....	12
3.1.3 hc:printf Functionality in ROCm Re-Enabled	12
Chapter 4 Known Limitations and Deprecations	13
4.1 Known Issues.....	13
4.1.1 Memory Access Fault Error While Running RCCL in Docker Container	13
4.1.2 Workaround for Tracer Library Fails to Load on RHEL.....	13
4.2 Deprecated Features.....	14
4.2.1 ROCm OpenCL Driver Compilation Services	14
4.2.2 Peer-to-Peer Bridge Driver for PeerDirect	14
Chapter 5 Deploying ROCm.....	14
5.1 Ubuntu	15
5.1.1 Installing a ROCm Package from a Debian Repository	15
5.1.2 Uninstalling ROCm Packages from Ubuntu.....	16
5.1.3 Installing Development Packages for Cross Compilation	17

5.1.4	Using Debian-based ROCm with Upstream Kernel Drivers	17
5.2	CentOS/RHEL v7 (v7.6) Support	17
5.2.1	Preparing RHEL v7 (7.6) for Installation.....	18
5.2.2	Installing CentOS/RHEL (v7.6) for DKMS.....	18
5.3	ROCm Installation.....	19
5.3.1	Installing ROCm	19
5.3.2	Setting Permissions	19
5.3.3	Testing the ROCm Installation.....	20
5.3.4	Performing an OpenCL-only Installation of ROCm	20
5.3.5	Uninstalling ROCm from CentOS/RHEL v7.6.....	21
5.3.6	Installing Development Packages for Cross Compilation.....	21
5.3.7	Using ROCm with Upstream Kernel Drivers	21
5.3.8	ROCm Installation - Known Issues and Workarounds	21
5.4	Getting the ROCm Source Code	22
5.4.1	Installing the Repo	22
5.4.2	Downloading the ROCm Source Code	22
5.4.3	Building the ROCm Source Code	23
Chapter 6	Hardware and Software Support.....	23
6.1	Hardware Support	23
6.1.1	Supported Graphics Processing Units	23
6.1.2	Supported CPUs	25
6.2	Software Support.....	28

Revision History

Date	Revision	Description
October 2019	1.0	Initial preliminary release
November 2019	2.0	Revision for new features
November 2019	2.1	Updated deprecation and known issues

Chapter 1 Introduction

This document describes the features, fixed issues, and information about downloading and installing the ROCm software.

It also covers known issues and deprecated features in the ROCm v2.10 release.

1.1 What Is ROCm?

ROCm is designed to be a universal platform for gpu-accelerated computing. This modular design allows hardware vendors to build drivers that support the ROCm framework. ROCm is also designed to integrate multiple programming languages and makes it easy to add support for other languages.

ROCm is built from open source software. Subject to the applicable license, you can download the source code, modify and rebuild the ROCm components. To ensure you are downloading the correct source code versions, the ROCm repository provides a repo manifest file called *default.xml*.

Note: You can also clone the source code for individual ROCm components from the GitHub repositories.

1.1.1 ROCm Components

The following components for the ROCm platform are released and available for the v2.10 release:

- Drivers
- ToolChains
- Libraries
- Source Code

You can access the latest supported version of drivers, tools, libraries, and source code for the ROCm platform at the following location:

<https://github.com/RadeonOpenCompute/ROCm>

1.1.2 Supported Operating Systems

The ROCm v2.10.x platform is designed to support the following operating systems:

- SLES 15 SP1
- Ubuntu 16.04.6(Kernel 4.15) and 18.04.3(Kernel 5.0)
- CentOS 7.6 (Using devtoolset-7 runtime support)
- RHEL 7.6 (Using devtoolset-7 runtime support)

For details about deploying the ROCm v2.10.x on these operating systems, see the [Deploying ROCm](#) section later in the document.

1.1.3 Important ROCm Links

Access the following links for more information on:

- ROCm documentation, see <https://rocm-documentation.readthedocs.io/en/latest/index.html>
- ROCm binary structure, see <https://github.com/RadeonOpenCompute/ROCm/blob/master/README.md#rocm-binary-package-structure>
- Common ROCm installation issues, see https://rocm.github.io/install_issues.html
- Instructions to install PyTorch after ROCm is installed – https://rocm-documentation.readthedocs.io/en/latest/Deep_learning/Deep-learning.html#pytorch

Note: These instructions reference the rocm/pytorch:rocm2.9_ubuntu16.04_py2.7_pytorch image. However, you can substitute the Ubuntu 18.04 image listed at <https://hub.docker.com/r/rocm/pytorch/tags>

Chapter 2 What's New in Version 2.10

2.1 New Features

2.1.1 rocBLAS - Support for Complex GEMM

The rocBLAS library is a gpu-accelerated implementation of the standard Basic Linear Algebra Subroutines (BLAS). rocBLAS is designed to enable you to develop algorithms, including high performance computing, image analysis, and machine learning.

In the AMD ROCm release v2.10, support is extended to the General Matrix Multiply (GEMM) routine for multiple small matrices processed simultaneously for rocBLAS in AMD Radeon Instinct MI50. Both single and double precision, CGEMM and ZGEMM, are now supported in rocBLAS.

2.1.2 Support for SLES 15 SP1

In the AMD ROCm v2.10 release, support is added for SUSE Linux® Enterprise Server (SLES) 15 SP1. SLES is a modular operating system for both multimodal and traditional IT.

Note: The SUSE Linux® Enterprise Server is a licensed platform. Ensure you have registered and have a license key prior to installation. Use the following SUSE command line to apply your license:

```
SUSEConnect -r < Key>
```

2.1.2.1 SLES 15 SP1

The following section tells you how to perform an install and uninstall ROCm on SLES 15 SP 1.

Run the following commands once for a fresh install on the operating system:

```
sudo usermod -a -G video $LOGNAME
```

```
sudo usermod -a -G sudo $LOGNAME
```

```
sudo reboot
```

Installation

1. Install the "dkms" package.

```
sudo SUSEConnect --product PackageHub/15.1/x86_64  
sudo zypper install dkms
```

2. Add the ROCm repo.

```
sudo zypper clean --all  
sudo zypper addrepo --no-gpgcheck http://repo.radeon.com/rocm/zypper/ rocm  
sudo zypper ref  
zypper install rocm-dkms  
sudo zypper install rocm-dkms  
sudo reboot  
  
#Run the following command once  
cat <<EOF | sudo tee /etc/modprobe.d/10-unsupported-modules.conf  
  
allow_unsupported_modules 1  
  
EOF  
  
sudo modprobe amdgpu
```

3. Verify the ROCm installation.

Run `/opt/rocm/bin/rocm_info` and `/opt/rocm/ocl/bin/x86_64/clinfo` commands to list the GPUs and verify that the ROCm installation is successful.

Uninstallation

To uninstall, use the following command:

```
sudo zypper remove rocm-dkms rock-dkms  
  
#Ensure all other installed packages/components are removed
```

Note: Ensure all the content in the `/opt/rocm` directory is completely removed.

For installation and uninstallation on deploying ROCm on other operating systems, see [Deploying ROCm](#).

2.1.3 Code Marker Support for rocProfiler and rocTracer Libraries

Code markers provide the external correlation ID for the calling thread. This function indicates that the calling thread is entering and leaving an external API region.

- The rocProfiler library enables you to profile performance counters and derived metrics. This library supports GFX8/GFX9 and provides a hardware-specific low-level performance analysis interface for profiling of GPU compute applications. The profiling includes hardware performance counters with complex performance metrics.
- The rocTracer library provides a specific runtime profiler to trace API and asynchronous activity. The API provides functionality for registering the runtimes API callbacks and the asynchronous activity records pool support.
- rocTX provides a C API for code markup for performance profiling and supports annotation of code ranges and ASCII markers.

Chapter 3 Fixed Issues

3.1 Fixed Issues in the v2.10 Release

3.1.1 Running TensorFlow and PyTorch Frameworks Consecutively Results in the Memory Access Fault error

Issue: Running the TensorFlow and PyTorch frameworks in quick succession results in a Memory Access Fault error.

Resolution: This issue is resolved, and the error no longer appears.

3.1.2 Issue with the Docker Container Environment Variable Setting

Issue: Applications fail when the docker container is launched on the NUMA system without the ‘security-opt seccomp=unconfined’ setting.

Resolution: Configure the “--security-opt seccomp=unconfined” variable setting to avoid this issue.

3.1.3 hc:printf Functionality in ROCm Re-Enabled

Known issues with hc:printf have been addressed in ROCm v2.10. The hc:printf functionality has now been re-enabled on all supported distros.

Chapter 4 Known Limitations and Deprecations

4.1 Known Issues

4.1.1 Memory Access Fault Error While Running RCCL in Docker Container

Issue: The **Memory Access Fault** error appears while running ROCm Communication Collectives Library (RCCL) tests in the Docker container.

Resolution: While launching the Docker container to run tests related to RCCL, including PyTorch, increase the size limit for the shared memory (SHM) directory to 1 GB. To increase the size limit of the shared memory directory, enter:

```
--shm-size = 1G
```

By default, Docker uses only 16 MB of shared memory. Running a Docker container for RCCL requires you to resize the limit to 1 GB.

4.1.2 Workaround for Tracer Library Fails to Load on RHEL

Issue: When running `/opt/rocm/bin/rocprof --hip-trace <output filename>`, a warning message is printed to console: "Tool lib "/opt/rocm/roctracer/tool/libtracer_tool.so" failed to load", and no output file is generated, on systems with RHEL distro.

Resolution: You can use either of the following workarounds to fix the issue:

- Run `Idconfig`
'SUDO LDCONFIG'

or

- Configure `LD_LIBRARY_PATH`
'EXPORT LD_LIBRARY_PATH=\${LD_LIBRARY_PATH}:/opt/rocm/roctracer/lib'

4.2 Deprecated Features

The following features are deprecated in the AMD ROCm v2.10 release.

4.2.1 ROCm OpenCL Driver Compilation Services

The AMD ROCm-OpenCL-Driver Compilation Services is now deprecated. Users should migrate to ROCm-CompilerSupport, which provides more comprehensive functionality. The compiler support repository provides various lightning compiler-related services. It currently contains a single library, the Code Object Manager (Comgr) at lib/comgr.

The ROCm OpenCL Driver Compilation Services will no longer be actively maintained after the v2.10 release. We would encourage you to switch to the ROCm-CompilerSupport repository.

4.2.2 Peer-to-Peer Bridge Driver for PeerDirect

The Peer-to-Peer bridge driver for the PeerDirect feature still works in the current release, however, it is now included as part of the ROCK kernel driver. ROCmRDMA allows third-party kernel drivers to utilize DMA access to the GPU memory. It allows a direct path for data exchange (peer-to-peer) using the standard features of PCI Express.

Currently, ROCmRDMA provides the following benefits:

- Direct access to ROCm memory for 3rd party PCIe devices
- Support for PeerDirect(c) interface to offloads the CPU when dealing with ROCm memory for RDMA network stacks

Chapter 5 Deploying ROCm

AMD hosts both Debian and RPM repositories for the ROCm v2.10x packages.

The following directions show how to install ROCm on supported Debian-based systems such as Ubuntu 18.04.

Note: These directions may not work as written on unsupported Debian-based distributions. For example, newer versions of Ubuntu may not be compatible with the `rock-dkms` kernel driver. In this case, you can exclude the `rocm-dkms` and `rock-dkms` packages.

For more information on the ROCm binary structure, see <https://github.com/RadeonOpenCompute/ROCm/blob/master/README.md#rocm-binary-package-structure>

For information about upstream kernel drivers, see the [Using Debian-based ROCm with Upstream Kernel Drivers](#) section.

5.1 Ubuntu

5.1.1 Installing a ROCm Package from a Debian Repository

To install from a Debian Repository:

1. Run the following code to ensure that your system is up to date:

```
sudo apt update  
sudo apt dist-upgrade  
sudo apt install libnuma-dev  
sudo reboot
```

2. Add the ROCm apt repository.

For Debian-based systems like Ubuntu, configure the Debian ROCm repository as follows:

```
wget -q0 --  
http://repo.radeon.com/rocm/apt/debian/rocm.gpg.key | sudo apt-key add -echo 'deb  
[arch=amd64]  
http://repo.radeon.com/rocm/apt/debian/ xenial main' | sudo tee  
/etc/apt/sources.list.d/rocm.list
```

The gpg key may change; ensure it is updated when installing a new release. If the key signature verification fails while updating, re-add the key from the ROCm apt repository.

The current rocm.gpg.key is not available in a standard key ring distribution, but has the following `sha1sum` hash:

```
e85a40d1a43453fe37d63aa6899bc96e08f2817a rocm.gpg.key
```

3. Install the ROCm meta-package.

Update the appropriate repository list and install the `rocm-dkms` meta-package:

```
sudo apt update  
sudo apt install rocm-dkms
```

4. Set permissions.

To access the GPU, you must be a user in the video group. Ensure your user account is a

member of the video group prior to using ROCm. To identify the groups you are a member of, use the following command:

```
groups
```

5. To add your user to the video group, use the following command for the *sudo* password:

```
sudo usermod -a -G video $LOGNAME
```

6. By default, add any future users to the video group. Run the following command to add users to the video group:

```
echo 'ADD_EXTRA_GROUPS=1' | sudo tee -a /etc/adduser.conf
```

```
echo 'EXTRA_GROUPS=video' | sudo tee -a /etc/adduser.conf
```

7. Restart the system.
8. Test the basic ROCm installation.
9. After restarting the system, run the following commands to verify that the ROCm installation is successful. If you see your GPUs listed by both commands, the installation is considered successful.

```
/opt/rocm/bin/rocminfo
```

```
/opt/rocm/rocm/bin/x86_64/clinfo
```

Note: To run the ROCm programs more efficiently, add the ROCm binaries in your PATH.

```
echo 'export PATH=$PATH:/opt/rocm/bin:/opt/rocm/profiler/bin:/opt/rocm/rocm/bin/x86_64' |  
sudo tee -a /etc/profile.d/rocm.sh
```

If you have an installation issue, refer the FAQ at:

https://rocm.github.io/install_issues.html

5.1.2 Uninstalling ROCm Packages from Ubuntu

To uninstall the ROCm packages from **Ubuntu 16.04** or **Ubuntu 18.04**, run the following command:

```
sudo apt autoremove rocm-dkms rocm-dev rocm-utils
```


5.1.3 Installing Development Packages for Cross Compilation

It is recommended that you develop and test development packages on different systems. For example, some development or build systems may not have an AMD GPU installed. In this scenario, you must avoid installing the ROCk kernel driver on the development system. Instead, install the following development subset of packages:

```
sudo apt update
```

```
sudo apt install rocm-dev
```

Note: To execute ROCm enabled applications, you must install the full ROCm driver stack on your system.

5.1.4 Using Debian-based ROCm with Upstream Kernel Drivers

You can install the ROCm user-level software without installing the AMD's custom ROCk kernel driver. To use the upstream kernels, run the following commands instead of installing `rocm-dkms`:

```
sudo apt update
```

```
sudo apt install rocm-dev
```

```
echo 'SUBSYSTEM=="kfd", KERNEL=="kfd", TAG+="uaccess", GROUP="video"' | sudo tee  
/etc/udev/rules.d/70-kfd.rules
```

5.2 CentOS/RHEL v7 (v7.6) Support

This section describes how to install ROCm on supported RPM-based systems such as CentOS v7.6.

Note: The following instructions may not work on unsupported RPM-based distributions. For example, Fedora may not be compatible with the `rock-dkms` kernel driver. You can exclude the `rocm-dkms` and `rock-dkms` packages and use the upstream kernel driver instead.

Note: Although support for CentOS/RHEL v7 was added in ROCm v1.8, ROCm requires a special runtime environment provided by the RHEL Software Collections and additional `dkms` support packages to install and run correctly.

For more details, refer:

<https://github.com/RadeonOpenCompute/ROCm/blob/master/README.md#rocm-binary-package-structure>

5.2.1 Preparing RHEL v7 (7.6) for Installation

RHEL is a subscription-based operating system. You must enable the external repositories to install on the devtoolset-7 environment and the dkms support files.

Note: The following steps do not apply to the CentOS installation.

1. The subscription for RHEL must be enabled and attached to a pool ID. See the Obtaining an RHEL image and license page for instructions on registering your system with the RHEL subscription server and attaching to a pool id.
2. Enable the following repositories:

```
sudo subscription-manager repos --enable rhel-server-rhsccl-7-rpms  
sudo subscription-manager repos --enable rhel-7-server-optional-rpms  
sudo subscription-manager repos --enable rhel-7-server-extras-rpms
```

3. Enable additional repositories by downloading and installing the epel-release-latest-7 repository RPM:

```
sudo rpm -ivh
```

For more details, see

<https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm>

4. Install and set up Devtoolset-7.

To setup the Devtoolset-7 environment, follow the instructions on this page:

<https://www.softwarecollections.org/en/scls/rhsccl/devtoolset-7/>

Note: devtoolset-7 is a software collections package and is not supported by AMD.

5.2.2 Installing CentOS/RHEL (v7.6) for DKMS

Use the dkms tool to install the kernel drivers on CentOS/RHEL v7.6:

```
sudo yum install -y epel-release  
sudo yum install -y dkms kernel-headers-`uname -r` kernel-devel-`uname -r`
```

5.3 ROCm Installation

5.3.1 Installing ROCm

To install ROCm on your system, follow the instructions below:

1. Delete the previous versions of ROCm before installing the latest version.
2. Create a `/etc/yum.repos.d/rocm.repo` file with the following contents:

```
[ROCm]

name=ROCm

baseurl=http://repo.radeon.com/rocm/yum/rpm

enabled=1

gpgcheck=0
```

Note: The URL of the repository must point to the location of the repositories' repodata database.

3. Install ROCm components using the following command:

```
sudo yum install rocm-dkms
```

4. Restart the system.

The `rocm-dkms` component is installed and the `/dev/kfd` device is now available.

5.3.2 Setting Permissions

To configure permissions, following the instructions below:

1. Ensure that your user account is a member of the "video" or "wheel" group prior to using the ROCm driver. You can find which groups you are a member of with the following command:

```
groups
```

2. Add your user to the video (or wheel) group you will need the sudo password and can use the following command:

```
sudo usermod -a -G video $LOGNAME
```

Note: All future users must be added to the "video" group by default. To add the users to the group, run the following commands

```
echo 'ADD_EXTRA_GROUPS=1' | sudo tee -a /etc/adduser.conf  
echo 'EXTRA_GROUPS=video' | sudo tee -a /etc/adduser.conf
```

Note: The current release supports CentOS/RHEL v7.6. Before updating to the latest version of the operating system, delete the ROCm packages to avoid DKMS-related issues.

3. Restart the system.

5.3.3 Testing the ROCm Installation

After restarting the system, run the following commands to verify that the ROCm installation is successful. If you see your GPUs listed, you are good to go!

```
/opt/rocm/bin/rocminfo  
  
/opt/rocm/opencl/bin/x86_64/clinfo
```

Note: Add the ROCm binaries in your PATH for easy implementation of the ROCm programs.

```
echo 'export PATH=$PATH:/opt/rocm/bin:/opt/rocm/profiler/bin:/opt/rocm/opencl/bin/x86_64' | sudo  
tee -a /etc/profile.d/rocm.sh
```

For more information about installation issues, see:

https://rocm.github.io/install_issues.html

5.3.4 Performing an OpenCL-only Installation of ROCm

Some users may want to install a subset of the full ROCm installation. If you are trying to install on a system with a limited amount of storage space, or which will only run a small collection of known applications, you may want to install only the packages that are required to run OpenCL applications. To do that, you can run the following installation command **instead** of the command to install rocm-dkms.

```
sudo yum install rock-dkms rocm-opencl-devel
```

5.3.4.1 Compiling Applications Using HCC, HIP, and Other ROCm Software

To compile applications or samples, run the following command to use gcc-7.2 provided by the devtoolset-7 environment:

```
scl enable devtoolset-7 bash
```

5.3.5 Uninstalling ROCm from CentOS/RHEL v7.6

To uninstall the ROCm packages, run the following command:

```
sudo yum autoremove rocm-dkms rock-dkms
```

5.3.6 Installing Development Packages for Cross Compilation

You can develop and test ROCm packages on different systems. For example, some development or build systems may not have an AMD GPU installed. In this scenario, you can avoid installing the ROCm kernel driver on your development system. Instead, install the following development subset of packages:

```
sudo yum install rocm-dev
```

Note: To execute ROCm-enabled applications, you will require a system installed with the full ROCm driver stack.

5.3.7 Using ROCm with Upstream Kernel Drivers

You can install ROCm user-level software without installing AMD's custom ROCk kernel driver. To use the upstream kernel drivers, run the following commands

```
sudo yum install rocm-dev
```

```
echo 'SUBSYSTEM=="kfd", KERNEL=="kfd", TAG+="uaccess", GROUP="video" | sudo tee /etc/udev/rules.d/70-kfd.rules
```

Note: You can use this command instead of installing *rocm-dkms*.

5.3.8 ROCm Installation - Known Issues and Workarounds

5.3.8.1 Docker container environment variable setting

Issue: Applications fail when a Docker container is launched on a NUMA system without `--security-opt seccomp=unconfined`.

Resolution: Set `"--security-opt seccomp=unconfined"` to fix this issue.

5.3.8.2 Closed source components

The ROCm platform relies on some closed source components to provide functionalities like HSA image support. These components are only available through the ROCm repositories, and they may be deprecated or become open source components in the future. These components are made available in the following packages:

- *hsa-ext-rocr-dev*

5.4 Getting the ROCm Source Code

AMD ROCm is built from open source software. It is, therefore, possible to modify the various components of ROCm by downloading the source code and rebuilding the components. The source code for ROCm components can be cloned from each of the GitHub repositories using git. For easy access to download the correct versions of each of these tools, the ROCm repository contains a repo manifest file called `default.xml`. You can use this manifest file to download the source code for ROCm software.

5.4.1 Installing the Repo

The repo tool from Google® allows you to manage multiple git repositories simultaneously. Run the following commands to install the repo:

```
mkdir -p ~/bin/  
  
curl https://storage.googleapis.com/git-repo-downloads/repo > ~/bin/repo  
  
chmod a+x ~/bin/repo
```

Note: You can choose a different folder to install the repo into if you desire. `~/bin/` is used as an example.

5.4.2 Downloading the ROCm Source Code

The following example shows how to use the repo binary to download the ROCm source code. If you choose a directory other than `~/bin/` to install the repo, you must use that chosen directory in the code as shown below:

```
mkdir -p ~/ROCM/  
  
cd ~/ROCM/  
  
~/bin/repo init -u https://github.com/RadeonOpenCompute/ROCM.git -b roc-2.10.0
```

repo sync

Note: Using this sample code will cause the repo to download the open source code associated with this ROCm release. Ensure that you have ssh-keys configured on your machine for your GitHub ID prior to the download.

5.4.3 Building the ROCm Source Code

Each ROCm component repository contains directions for building that component. You can access the desired component for instructions to build the repository.

Chapter 6 Hardware and Software Support

6.1 Hardware Support

ROCm is focused on using AMD GPUs to accelerate computational tasks such as machine learning, engineering workloads, and scientific computing. In order to focus our development efforts on these domains of interest, ROCm supports the following targeted set of hardware configurations.

6.1.1 Supported Graphics Processing Units

As the AMD ROCm platform has a focus on specific computational domains, AMD offers official support for a selection of GPUs that are designed to offer good performance and price in these domains.

ROCm officially supports AMD GPUs that use the following chips:

- GFX8 GPUs
 - "Fiji" chips, such as on the AMD Radeon R9 Fury X and Radeon Instinct MI8
 - "Polaris 10" chips, such as on the AMD Radeon RX 580 and Radeon Instinct MI6
- GFX9 GPUs
 - "Vega 10" chips, such as on the AMD Radeon RX Vega 64 and Radeon Instinct MI25
 - "Vega 7nm" chips, such as on the Radeon Instinct MI50, Radeon Instinct MI60 or AMD Radeon VII

ROCm is a collection of software ranging from drivers and runtimes to libraries and developer tools. Some of this software may work with more GPUs than the "officially supported" list above, though AMD does not make any official claims of support for these devices on the ROCm software platform. The following list of GPUs is enabled in the ROCm software. However, full support is not guaranteed:

- GFX8 GPUs
 - "Polaris 11" chips, such as on the AMD Radeon RX 570 and Radeon Pro WX 4100
 - "Polaris 12" chips, such as on the AMD Radeon RX 550 and Radeon RX 540
- GFX7 GPUs
 - "Hawaii" chips, such as the AMD Radeon R9 390X and FirePro W9100

As described in the next section, GFX8 GPUs require PCI Express 3.0 (PCIe 3.0) with support for PCIe atomics. This requires both CPU and motherboard support. GFX9 GPUs require PCIe 3.0 with support for PCIe atomics by default, but they can operate in most cases without this capability.

The integrated GPUs in AMD APU are not officially supported targets for ROCm. As described below, "Carrizo", "Bristol Ridge", and "Raven Ridge" APUs are enabled in AMD upstream drivers and the ROCm OpenCL runtime. However, they are not enabled in AMD HCC or HIP runtimes, and may not work due to motherboard or OEM hardware limitations. Note, they are not yet officially supported targets for ROCm.

6.1.1.1 GFX8 GPUs

ROCm offers support for the following microprocessors from AMD's "gfx8" generation of GPUs.

Note: The GPUs require a host CPU and platform with PCIe 3.0 with support for PCIe atomics.

GFX8 GPUs			
Fiji (AMD)	Polaris 10 (AMD)	Polaris 11 (AMD)	Polaris 12 (Lexa) (AMD)
<ul style="list-style-type: none"> • Radeon R9 Fury • Radeon R9 Nano • Radeon R9 Fury X • Radeon Pro Duo (Fiji) • FirePro S9300 X2 • Radeon Instinct MI8 	<ul style="list-style-type: none"> • Radeon RX 470 • Radeon RX 480 • Radeon RX 570 • Radeon RX 580 • Radeon Pro Duo (Polaris) • Radeon Pro WX 5100 • Radeon Pro 	<ul style="list-style-type: none"> • Radeon RX 460 • Radeon RX 560 • Radeon Pro WX 4100 	<ul style="list-style-type: none"> • Radeon RX 540 • Radeon RX 550 • Radeon Pro WX 2100 • Radeon Pro WX 3100

	<ul style="list-style-type: none"> WX 7100 Radeon Instinct MI6 		
--	--	--	--

6.1.1.2 GFX9 GPUs

ROCm offers support for two chips from AMD’s most recent “gfx9” generation of GPUs.

GFX9 GPUs	
Vega 10 (AMD)	Vega 7nm (AMD)
<ul style="list-style-type: none"> Radeon RX Vega 56 Radeon RX Vega 64 Radeon Vega Frontier Edition Radeon Pro WX 8200 Radeon Pro WX 9100 Radeon Pro V340 Radeon Pro V340 MxGPU Radeon Instinct MI25 	<ul style="list-style-type: none"> Radeon VII Radeon Instinct MI50 Radeon Instinct MI60
Note: ROCm does not support Radeon Pro SSG.	

6.1.2 Supported CPUs

As described above, GFX8 GPUs require PCIe 3.0 with PCIe atomics to run ROCm. In particular, the CPU and every active PCIe point between the CPU and GPU require support for PCIe 3.0 and PCIe atomics. The CPU root must indicate PCIe AtomicOp Completion capabilities and any intermediate switch must indicate PCIe AtomicOp Routing capabilities.

The current CPUs which support PCIe Gen3 + PCIe Atomics are:

- AMD Ryzen CPUs
- The CPUs in AMD Ryzen APUs
- AMD Ryzen Threadripper CPUs
- AMD EPYC CPUs
- Intel Xeon E7 v3 or newer CPUs
- Intel Xeon E5 v3 or newer CPUs
- Intel Xeon E3 v3 or newer CPUs
- Intel Core i7 v4, Core i5 v4, Core i3 v4 or newer CPUs (i.e. Haswell family or newer)
- Some Ivy Bridge-E systems

Beginning with ROCm 1.8, GFX9 GPUs (such as Vega 10) no longer require PCIe atomics. We have similarly made more options available for many PCIe lanes. GFX9 GPUs can now be run on CPUs without PCIe atomics and on older PCIe generations, such as PCIe 2.0. This is not supported on GPUs below GFX9, e.g. GFX8 cards in the Fiji and Polaris families.

If you are using any PCIe switches in your system, please note that PCIe Atomics are only supported on some switches, such as Broadcom PLX. When you install your GPUs, make sure you install them in a PCIe 3.0 x16, x8, x4, or x1 slot attached either directly to the CPU's Root I/O controller or via a PCIe switch directly attached to the CPU's Root I/O controller.

In our experience, many issues stem from trying to use consumer motherboards which provide physical x16 connectors that are electrically connected as e.g. PCIe 2.0 x4, PCIe slots connected via the Southbridge PCIe I/O controller, or PCIe slots connected through a PCIe switch that does not support PCIe atomics.

If you attempt to run ROCm on a system without proper PCIe atomic support, you may see an error in the kernel log (`dmesg`):

```
kfd: skipped device 1002:7300, PCI rejects atomics
```

Experimental support for our Hawaii (GFX7) GPUs (Radeon R9 290, R9 390, FirePro W9100, S9150, S9170) does not require or take advantage of PCIe Atomics. However, AMD recommends that you use a CPU from the list provided above for compatibility purposes.

6.1.2.1 Not supported or limited support under ROCm

6.1.2.1.1 Limited support

- ROCm 2.10.x should support PCIe 2.0 enabled CPUs such as the AMD Opteron, Phenom, Phenom II, Athlon, Athlon X2, Athlon II and older Intel Xeon and Intel Core Architecture and Pentium CPUs. However, we have done very limited testing on these configurations, since our test farm has been catering to CPUs listed above. This is where we need community support. *If you find problems on such setups, please report these issues.*
- Thunderbolt 1, 2, and 3 enabled breakout boxes should now be able to work with ROCm. Thunderbolt 1 and 2 are PCIe 2.0 based, and thus are only supported with GPUs that do not require PCIe 3.0 atomics (e.g. Vega 10). However, we have done no testing on this configuration and would need community support due to limited access to this type of equipment.
- AMD "Carrizo" and "Bristol Ridge" APUs are enabled to run OpenCL, but do not yet support HCC, HIP, or our libraries built on top of these compilers and runtimes.
 - As of ROCm 2.1, "Carrizo" and "Bristol Ridge" require the use of upstream kernel drivers.
 - In addition, various "Carrizo" and "Bristol Ridge" platforms may not work due to OEM and ODM choices when it comes to key configurations parameters such as inclusion of the required CRAT tables and IOMMU configuration parameters in the system BIOS.

- Before purchasing such a system for ROCm, please verify that the BIOS provides an option for enabling IOMMUv2 and that the system BIOS properly exposes the correct CRAT table. Inquire with your vendor about the latter.
- AMD "Raven Ridge" APU's are enabled to run OpenCL, but do not yet support HCC, HIP, or our libraries built on top of these compilers and runtimes.
 - As of ROCm 2.1, "Raven Ridge" requires the use of upstream kernel drivers.
 - In addition, various "Raven Ridge" platforms may not work due to OEM and ODM choices when it comes to key configuration parameters such as inclusion of the required CRAT tables and IOMMU configuration parameters in the system BIOS.
 - Before purchasing such a system for ROCm, please verify that the BIOS provides an option for enabling IOMMUv2 and that the system BIOS properly exposes the correct CRAT table. Inquire with your vendor about the latter.

6.1.2.1.2 Not supported

- "Tonga", "Iceland", "Vega M", and "Vega 12" GPUs are not supported.
- AMD does not support GFX8-class GPUs (Fiji, Polaris, etc.) on CPUs that do not have PCIe3.0 with PCIe atomics.
 - AMD Carrizo and Kaveri APU's as hosts for such GPUs are not supported
 - Thunderbolt 1 and 2 enabled GPUs are not supported by GFX8 GPUs on ROCm. Thunderbolt 1 & 2 are based on PCIe 2.0.

In the default ROCm configuration, GFX8 and GFX9 GPUs require PCI Express 3.0 with PCIe atomics. The ROCm platform leverages these advanced capabilities to allow features such as user-level submission of work from the host to the GPU. This includes PCIe atomic Fetch and Add, Compare and Swap, Unconditional Swap, and AtomicOp Completion.

Current CPUs which support PCIe 3.0 + PCIe Atomics:

AMD	INTEL
Ryzen CPUs (Family 17h Model 01h-0Fh) <ul style="list-style-type: none"> • Ryzen 3 1300X • Ryzen 3 2300X • Ryzen 5 1600X • Ryzen 5 2600X • Ryzen 7 1800X • Ryzen 7 2700X 	Intel Core i3, i5, and i7 CPUs from Haswell and beyond. This includes: <ul style="list-style-type: none"> • Haswell CPUs such as the Core i7 4790K • Broadwell CPUs such as the Core i7 5775C • Skylake CPUs such as the Core i7 6700K • Kaby Lake CPUs such as the Core i7 7740X

	<ul style="list-style-type: none"> • Coffee Lake CPUs such as the Core i7 8700K • Xeon CPUs from “v3” and newer • Some models of “Ivy Bridge-E” processors
Ryzen APUs (Family 17h Model 10h-1Fh – previously code-named Raven Ridge) such as: <ul style="list-style-type: none"> • Athlon 200GE • Ryzen 5 2400G <p>Note: The integrated GPU in these devices is not guaranteed to work with ROCm.</p>	
Ryzen Threadripper Workstation CPUs (Family 17h Model 01h-0Fh) such as: <ul style="list-style-type: none"> • Ryzen Threadripper 1950X • Ryzen Threadripper 2990WX 	
EPYC Server CPUs (Family 17h Model 01h-0Fh) such as: <ul style="list-style-type: none"> • Epyc 7551P • Epyc 7601 	

6.2 Software Support

As of AMD ROCm v1.9.0, the ROCm user-level software is compatible with the AMD drivers in certain upstream Linux kernels. You have the following options:

- Use the ROCk kernel driver that is a part of AMD’s ROCm repositories
- or
- Use the upstream driver and only install ROCm user-level utilities from AMD’s ROCm repositories

The releases of the upstream Linux kernel support the following GPUs in ROCm:

- Fiji, Polaris 10, Polaris 11
- Fiji, Polaris 10, Polaris 11, Vega10
- Fiji, Polaris 10, Polaris 11, Vega10, Vega 7nm

6.2.1.1 Supported Products

- CUDA v8